

SE 4050 – Deep Learning

4th Year second semester



Lab Assignment 3

Lasal Sandeepa Hettiarachchi

IT19132310

B.Sc. (Hons) in Information Technology Specializing in Software Engineering

Submitted to

Department of Computer Science and Software Engineering

Sri Lanka Institute of Information Technology

Sri Lanka

TABLE OF CONTENTS

Table of Contents	ii
1. Model-based algorithms (Value iteration approach)	3
1.1 valueIterationAgents.py.....	3
1.2 analysis.py	4
2. Model-Free algorithms (Q-learning)	6
2.1 qlearningAgents.py	6
3. Autograder tool results	12
3.1 Question 1	12
3.2 Question 2	12
3.3 Question 3	13
3.4 Question 4	13
3.5 Question 5	13
3.6 Question 6	14
3.7 Question 7	14
4. Appendix	15

1. MODEL-BASED ALGORITHMS (VALUE ITERATION APPROACH)

Model-based algorithms require a model of the environment in that it incorporates the knowledge of the transition process from one state to another.

1.1 valueIterationAgents.py

```
2  # Write value iteration code here
3      """ YOUR CODE HERE """
4
5  for _ in range(iterations):
6      current_values = self.values.copy() #Getting a copy of the
7      states = self.mdp.getStates()
8      for state in states: # iterating through each state
9          if self.mdp.isTerminal(state):
10             continue
11             # get value for best possible action for changing state
12             actions = self.mdp.getPossibleActions(state)
13             val_arr = []
14             for a in actions:
15                 val_arr.append(self.getQValue(state, a))
16             best_value = max(val_arr)
17             current_values[state] = best_value
18
19     self.values = current_values


def computeQValueFromValues(self, state, action):
    """
    Compute the Q-value of action in state from the
    value function stored in self.values.
    """
    """ YOUR CODE HERE """
    qValue = 0

    # go through every possible outcome of the action
    t_state_probs = self.mdp.getTransitionStatesAndProbs(state, action)
    for nextState, probability in t_state_probs:

        # add reward & future reward (=V) * probability of the outcome
```

```

        reward = self.mdp.getReward(state, action, nextState)
        discount = self.discount
        val = self.values[nextState]
        qValue = qValue + probability * (reward + discount * val)

    return qValue
    util.raiseNotDefined()

def computeActionFromValues(self, state):
    """
    The policy is the best action in the given state
    according to the values currently stored in self.values.

    You may break ties any way you see fit. Note that if
    there are no legal actions, which is the case at the
    terminal state, you should return None.
    """
    """ YOUR CODE HERE """
    # retrieve the best possible action for the state
    policies = util.Counter()
    pos_actions = self.mdp.getPossibleActions(state)
    for action in pos_actions:

        # how good is an action = q-value (which considers all possible outcomes)
        policies[action] = self.getQValue(state, action)

    # return the best action, e.g. 'north'
    return policies.argmax()
    util.raiseNotDefined()

```

1.2 analysis.py

```

2 def question2():
3     answerDiscount = 0.9
4     answerNoise = 0
5     return answerDiscount, answerNoise
6
7 def question3a():

```

```

8     answerDiscount = 0.1
9     answerNoise = 0
10    answerLivingReward = 0.7 # give a negative penalty for moving
11    return answerDiscount, answerNoise, answerLivingReward
12
13    def question3b():
14        answerDiscount = 0.1
15        answerNoise = 0.1
16        answerLivingReward = 0.7
17        return answerDiscount, answerNoise, answerLivingReward
18
19    def question3c():
20        answerDiscount = 0.5
21        answerNoise = 0
22        answerLivingReward = 0.5
23        return answerDiscount, answerNoise, answerLivingReward
24
25    def question3d():
26        answerDiscount = 0.9
27        answerNoise = 0.1
28        answerLivingReward = 0.2
29        return answerDiscount, answerNoise, answerLivingReward
30
31    def question3e():
32        answerDiscount = 0
33        answerNoise = 0
34        answerLivingReward = 0
35        return answerDiscount, answerNoise, answerLivingReward
36
37    def question6():
38        answerEpsilon = None
39        answerLearningRate = None
40        # return answerEpsilon, answerLearningRate
41        # If not possible, return 'NOT POSSIBLE'
42        return 'NOT POSSIBLE'
43
44    if __name__ == '__main__':
45        print 'Answers to analysis questions:'
46        import analysis
47        for q in [q for q in dir(analysis) if q.startswith('question')]:

```

```

48     response = getattr(analysis, q)()
49     print ' Question %s:\t%s' % (q, str(response))

```

2. MODEL-FREE ALGORITHMS (Q-LEARNING)

In reinforcement learning, a model-free algorithm is an algorithm which does not use the transition probability distribution associated

2.1 qlearningAgents.py

```

3  # qlearningAgents.py
4  # -----
5  # Licensing Information: You are free to use or extend these projects for
6  # educational purposes provided that (1) you do not distribute or publish
7  # solutions, (2) you retain this notice, and (3) you provide clear
8  # attribution to UC Berkeley, including a link to http://ai.berkeley.edu.
9  #
10 # Attribution Information: The Pacman AI projects were developed at UC Berkeley.
11 # The core projects and autograders were primarily created by John DeNero
12 # (denero@cs.berkeley.edu) and Dan Klein (klein@cs.berkeley.edu).
13 # Student side autograding was added by Brad Miller, Nick Hay, and
14 # Pieter Abbeel (pabbeel@cs.berkeley.edu).
15
16 from game import *
17 from learningAgents import ReinforcementAgent
18 from featureExtractors import *
19
20 import random,util,math
21
22 class QLearningAgent(ReinforcementAgent):
23     """
24     Q-Learning Agent
25
26     Functions you should fill in:
27     - computeValueFromQValues
28     - computeActionFromQValues
29     - getQValue
30     - getAction
31     - update
32
33     Instance variables you have access to

```

```

34     - self.epsilon (exploration prob)
35     - self.alpha (learning rate)
36     - self.discount (discount rate)
37
38     Functions you should use
39     - self.getLegalActions(state)
40         which returns legal actions for a state
41     """
42     def __init__(self, **args):
43         "You can initialize Q-values here..."
44         ReinforcementAgent.__init__(self, **args)
45
46         """ YOUR CODE HERE """
47         self.values = util.Counter()
48
49     def getQValue(self, state, action):
50         """
51         Returns Q(state,action)
52         Should return 0.0 if we have never seen a state
53         or the Q node value otherwise
54         """
55         """ YOUR CODE HERE """
56         return self.values[(state, action)]
57         util.raiseNotDefined()
58
59     def computeValueFromQValues(self, state):
60         """
61         Returns max_action Q(state,action)
62         where the max is over legal actions. Note that if
63         there are no legal actions, which is the case at the
64         terminal state, you should return a value of 0.0.
65         """
66         """ YOUR CODE HERE """
67         maxQ = float('-inf')
68         for action in self.getLegalActions(state):
69             maxQ = max(maxQ, self.getQValue(state, action))
70         return maxQ if maxQ != float('-inf') else 0.0
71         util.raiseNotDefined()
72
73     def computeActionFromQValues(self, state):

```

```

74     """
75     Compute the best action to take in a state. Note that if there
76     are no legal actions, which is the case at the terminal state,
77     you should return None.
78     """
79     """ YOUR CODE HERE """
80     if len(self.getLegalActions(state)) == 0:
81         return None
82
83     bestQ = self.computeValueFromQValues(state)
84     bestActions = []
85     for action in self.getLegalActions(state):
86         if bestQ == self.getQValue(state, action):
87             bestActions.append(action)
88
89     return random.choice(bestActions)
90     util.raiseNotDefined()
91
92     def getAction(self, state):
93         """
94         Compute the action to take in the current state. With
95         probability self.epsilon, we should take a random action and
96         take the best policy action otherwise. Note that if there are
97         no legal actions, which is the case at the terminal state, you
98         should choose None as the action.
99
100         HINT: You might want to use util.flipCoin(prob)
101         HINT: To pick randomly from a list, use random.choice(list)
102         """
103         # Pick Action
104         legalActions = self.getLegalActions(state)
105         action = None
106         """ YOUR CODE HERE """
107         if util.flipCoin(self.epsilon):
108             action = random.choice(legalActions)
109         else:
110             action = self.computeActionFromQValues(state)
111
112         return action
113         util.raiseNotDefined()

```



```

114
115 def update(self, state, action, nextState, reward):
116     """
117     The parent class calls this to observe a
118     state = action => nextState and reward transition.
119     You should do your Q-Value update here
120
121     NOTE: You should never call this function,
122     it will be called on your behalf
123     """
124     """ YOUR CODE HERE """
125     oldValue = self.values[(state, action)]
126     newValue = reward + (self.discount * self.computeValueFromQValues(nextState))
127
128     self.values[(state, action)] = (1 - self.alpha) * oldValue + self.alpha * newValue
129     #util.raiseNotDefined()
130
131 def getPolicy(self, state):
132     return self.computeActionFromQValues(state)
133
134 def getValue(self, state):
135     return self.computeValueFromQValues(state)
136
137 class PacmanQAgent(QLearningAgent):
138     "Exactly the same as QLearningAgent, but with different default parameters"
139
140     def __init__(self, epsilon=0.05,gamma=0.8,alpha=0.2, numTraining=0, **args):
141         """
142         These default parameters can be changed from the pacman.py command line.
143         For example, to change the exploration rate, try:
144             python pacman.py -p PacmanQLearningAgent -a epsilon=0.1
145
146         alpha    - learning rate
147         epsilon  - exploration rate
148         gamma    - discount factor
149         numTraining - number of training episodes, i.e. no learning after these many episodes
150         """
151         args['epsilon'] = epsilon
152         args['gamma'] = gamma
153         args['alpha'] = alpha

```

```

154     args['numTraining'] = numTraining
155     self.index = 0 # This is always Pacman
156     QLearningAgent.__init__(self, **args)
157
158     def getAction(self, state):
159         """
160         Simply calls the getAction method of QLearningAgent and then
161         informs parent of action for Pacman. Do not change or remove this
162         method.
163         """
164         action = QLearningAgent.getAction(self,state)
165         self.doAction(state,action)
166         return action
167
168     class ApproximateQAgent(PacmanQAgent):
169         """
170         ApproximateQLearningAgent
171
172         You should only have to overwrite getQValue
173         and update. All other QLearningAgent functions
174         should work as is.
175         """
176         def __init__(self, extractor='IdentityExtractor', **args):
177             self.featExtractor = util.lookup(extractor, globals())()
178             PacmanQAgent.__init__(self, **args)
179             self.weights = util.Counter()
180
181         def getWeights(self):
182             return self.weights
183
184         def getQValue(self, state, action):
185             """
186             Should return Q(state,action) = w * featureVector
187             where * is the dotProduct operator
188             """
189             """ YOUR CODE HERE """
190             features = self.featExtractor.getFeatures(state, action)
191
192             sum = 0
193             for feature, value in features.iteritems():

```

```

194         sum += self.weights[feature] * value
195     return
196     util.raiseNotDefined()
197
198     def update(self, state, action, nextState, reward):
199         """
200         Should update your weights based on transition
201         """
202         """ YOUR CODE HERE """
203         newValue = reward + self.discount * self.computeValueFromQValues(nextState)
204         oldValue = self.getQValue(state, action)
205         difference = newValue - oldValue
206
207         features = self.feetExtractor.getFeatures(state, action)
208         for feature, value in features.iteritems():
209             self.weights[feature] += self.alpha * difference * features[feature]
210             #util.raiseNotDefined()
211
212     def final(self, state):
213         "Called at the end of each game."
214         # call the super-class final method
215         PacmanQAgent.final(self, state)
216
217         # did we finish training?
218         if self.episodesSoFar == self.numTraining:
219             # you might want to print your weights here for debugging
220             """ YOUR CODE HERE """
221             pass
222

```

3. AUTOGRADER TOOL RESULTS

3.1 Question 1

```
• (py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % python2 autograder.py -q q1

Starting on 10-31 at 4:37:35

Question q1
=====

*** PASS: test_cases/q1/1-tinygrid.test
*** PASS: test_cases/q1/2-tinygrid-noisy.test
*** PASS: test_cases/q1/3-bridge.test
*** PASS: test_cases/q1/4-discountgrid.test

### Question q1: 6/6 ###

Finished at 4:37:35

Provisional grades
=====
Question q1: 6/6
-----
Total: 6/6

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.

○ (py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % █
```

3.2 Question 2

```
• (py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % python2 autograder.py -q q2

Starting on 10-31 at 4:38:19

Question q2
=====

*** PASS: test_cases/q2/1-bridge-grid.test

### Question q2: 1/1 ###

Finished at 4:38:19

Provisional grades
=====
Question q2: 1/1
-----
Total: 1/1

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.
```

3.3 Question 3

```
(py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % python2 autograder.py -q q3

Starting on 10-31 at 4:38:54

Question q3
=====

*** PASS: test_cases/q3/1-question-3.1.test
*** PASS: test_cases/q3/2-question-3.2.test
*** PASS: test_cases/q3/3-question-3.3.test
*** PASS: test_cases/q3/4-question-3.4.test
*** PASS: test_cases/q3/5-question-3.5.test

### Question q3: 5/5 ###

Finished at 4:38:55

Provisional grades
=====
Question q3: 5/5
-----
Total: 5/5

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.
```

3.4 Question 4

```
(py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % python2 autograder.py -q q4

Starting on 10-31 at 4:39:55

Question q4
=====

*** PASS: test_cases/q4/1-tinygrid.test
*** PASS: test_cases/q4/2-tinygrid-noisy.test
*** PASS: test_cases/q4/3-bridge.test
*** PASS: test_cases/q4/4-discountgrid.test

### Question q4: 5/5 ###

Finished at 4:39:55

Provisional grades
=====
Question q4: 5/5
-----
Total: 5/5

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.
```

3.5 Question 5

```
(py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % python2 autograder.py -q q5

Starting on 10-31 at 4:40:23

Question q5
=====

*** PASS: test_cases/q5/1-tinygrid.test
*** PASS: test_cases/q5/2-tinygrid-noisy.test
*** PASS: test_cases/q5/3-bridge.test
*** PASS: test_cases/q5/4-discountgrid.test

### Question q5: 3/3 ###

Finished at 4:40:25

Provisional grades
=====
Question q5: 3/3
-----
Total: 3/3

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.
```

3.6 Question 6

```
(py2_venv) (base) lasalhettiarachchi@Lasals-MacBook-Pro reinforcement % python2 autograder.py -q q6

Starting on 10-31 at 4:40:58

Question q6
=====

*** PASS: test_cases/q6/grade-agent.test

### Question q6: 1/1 ###

Finished at 4:40:58

Provisional grades
=====
Question q6: 1/1
=====
Total: 1/1

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.
```

3.7 Question 7

[illegible]

4. APPENDIX

